Stand-Alone Self-Attention in Vision Models

Prajit Ramachandran* **Niki Parmar*** **Ashish Vaswani***

Irwan Bello **Anselm Levskaya†** **Jonathon Shlens**

Google Research, Brain Team
{prajit, nikip, avaswani}@google.com

Attention is all you need!

Aim

Developing a fully self-attention model for object detection and recognition using local stand-alone self attention blocks.

- Fewer parameters? (yes)
- Faster runtime? (yes)
- Better accuracy? (yes)

under certain circumstances

have successful examples in language modelling and generation.

Contribution: Spatial Self Attention Block

A stand-alone self-attention layer that can be used to replace spatial convolutions and building a fully attentional model.

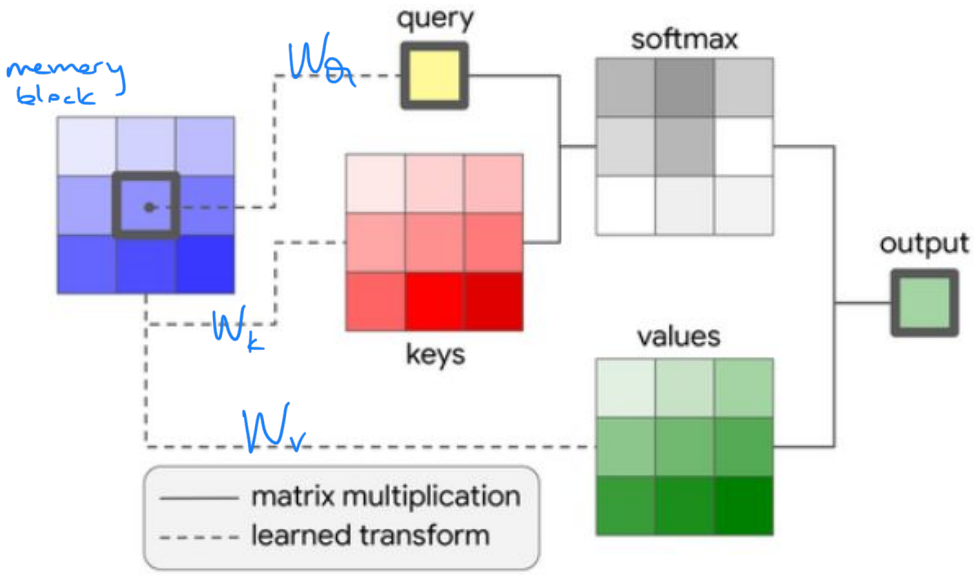

Relative Distance Computation

Using the relative distance to query pixel, spatial-relative attention is:

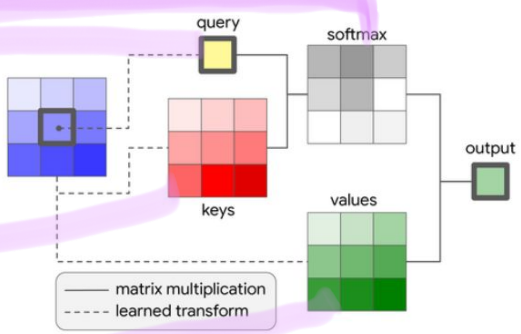

$$\text{out} = \sum \text{softmax}\left(q \times \text{keys} + q \times \underset{\text{pos.}}{\text{relative}}\right) \times \text{values}$$

-1, -1	-1, 0	-1, 1	-1, 2
0, -1	0, 0	0, 1	0, 2
1, -1	1, 0	1, 1	1, 2
2, -1	2, 0	2, 1	2, 2

Fully Attentional Vision Models

ResNet: – core mechanism:

1x1 conv → downsample

3x3 spatial conv

1x1 conv → upsample

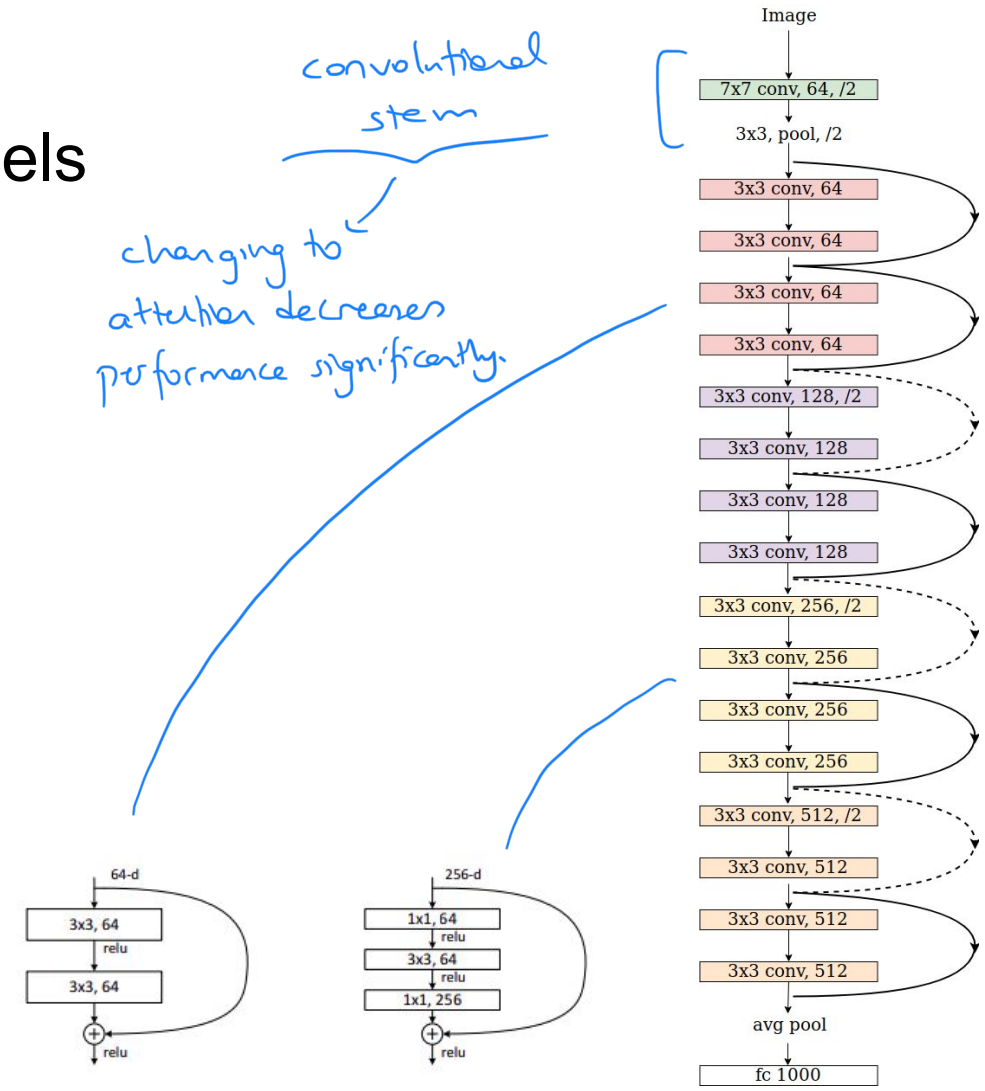

Experiments: ResNet on ImageNet

	ResNet-26			ResNet-38			ResNet-50		
	FLOPS (B)	Params (M)	Acc. (%)	FLOPS (B)	Params (M)	Acc. (%)	FLOPS (B)	Params (M)	Acc. (%)
Baseline	4.7	13.7	74.5	6.5	19.6	76.2	8.2	25.6	76.9
Conv-stem + Attention	4.5	10.3	**75.8**	5.7	14.1	**77.1**	7.0	18.0	**77.4**
Full Attention	4.7	10.3	74.8	6.0	14.1	76.9	7.2	18.0	**77.6**

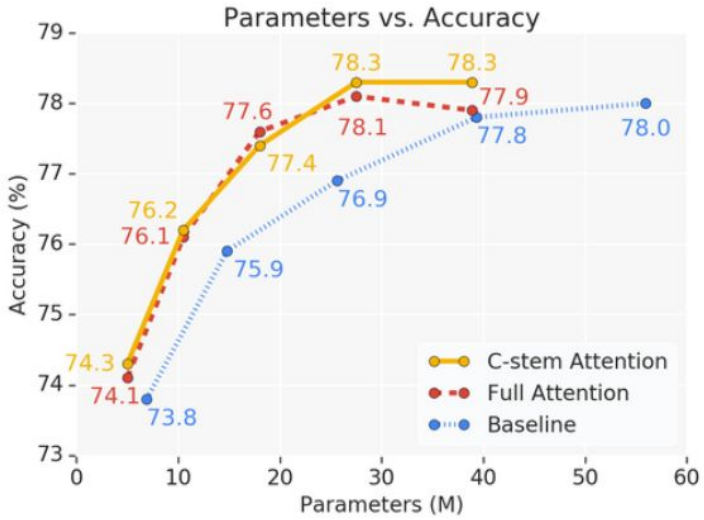

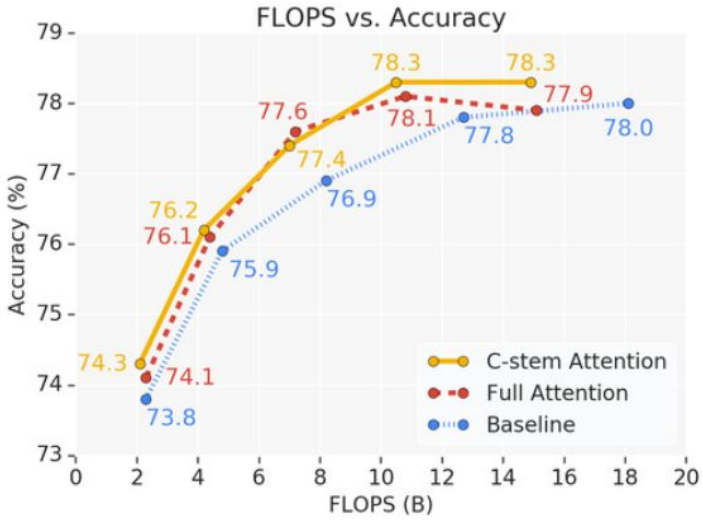

Which components are important in attention?

All these info is based on models with convolutional stem

- Spatial extent, k is 11 (improvement 3 -------> 11)
- Relative position encodings perform 2% better than absolute encodings.
-